

Resizing hash tables

- Fine so far, but will eventually be bad
- If $n = 100$ and size of hash table is 5, then will be (at best) 20 elements per index = $n/5$. Not good - as n increases, number of elements per bucket will be only a constant factor times n , so will eventually be $O(n)$ lookup again.
- Solution: resize hash table when there are too many elements in the table.
- Load factor: number of elements in table / number of table slots.
 - Typically want to resize hash table when load factor is around 0.75, but can vary by implementation.

Resizing hash tables (cont.)

- How can we resize efficiently? Do we need to rehash every single element? Would take $O(n)$ time (bad)!
- Answer: doesn't matter because we resize so infrequently.
- Insertion will be **amortized** $O(1)$ time:
 - Only resize once the load factor is 0.75
 - When we resize, make backing array twice as large
 - So, if we insert n elements, will only need to resize the array something like $\log(n)$ times - but each of them will only require $n/2^k$ rehashes, the sum of which is $O(n)$:
$$\sum_{k=0}^n \frac{n}{2^k} = n \sum_{k=0}^n \frac{1}{2^k} = 2n$$
- So, the total runtime impact of rehashes is $O(2\log(n)) = O(\log(n))$. However, we are already adding n elements, so the total amount of runtime for inserting n elements is $O(\log(n)) + O(n) = O(n)$.